

assignment08

April 14, 2021

```
[ ]: import numpy as np
import matplotlib.pyplot as plt
from sklearn.pipeline import Pipeline
from sklearn.metrics import mean_squared_error
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import mean_squared_error
```

```
[ ]: # Hidden parts create data x and y

x = np.linspace(-3,3, 30).reshape(-1, 1)
np.random.seed(42)
y_clean= x * -0.8545 + x**2*0.35+ x**3*0.25
y = y_clean + np.random.normal(np.zeros_like(x),1)

plt.plot(x, y, "x")
plt.xlabel("x")
plt.ylabel("y")
```

0.1 Regularisation - Ridge Regression

1. Create a polynomial input X_{poly} of degree 10 (without a bias term) from the sampling locations x
2. For $\alpha = 0, 1, 5, 10$ fit a Ridge Regressor to X_{poly} and y
3. Plot the predictions of the respective models over x and create a legend to distinguish them

You can check X_{poly} in the cell below.

```
[ ]: from sklearn.linear_model import Ridge

plt.plot(x, y, "x")
plt.xlabel("x")
plt.ylabel("y")

# YOUR CODE HERE
raise NotImplementedError()
```

```
[ ]: # Hidden part checks X_poly to avoid simple mistakes
```

```
if is_correct
    print("X_poly is correct!")
else:
    print("X_poly is not correct!")
```

0.2 Regularisation - Decision Tree

Tune the following parameters by hand to get a feel for their effect: * max_depth: Max depth of the tree * min_samples_split: The minimum number of samples required to split an internal node * min_samples_leaf: The minimum number of samples required to be at a leaf node.

The goal is to get a reasonable regressor. The original settings are below if you want to reset your try:

```
regressor = DecisionTreeRegressor(max_depth=None, min_samples_split=2, min_samples_leaf=1, random_state=0)
```

```
[ ]: from sklearn.tree import DecisionTreeRegressor
plt.plot(x, y, "x")
plt.xlabel("x")
plt.ylabel("y")
regressor = DecisionTreeRegressor(max_depth=None, min_samples_split=2,
    ↪min_samples_leaf=1, random_state=0)
regressor.fit(x,y)
y_pred = regressor.predict(x)
plt.plot(x, y_pred, label="alpha=")

# YOUR CODE HERE
raise NotImplementedError()
```

0.3 Cross Validation

To get to know cross validation you should see it in comparison to “traditional” train / test splitting. Therefore, implement the following steps:

- For a random seed from 1 to 5 split the X_poly and y into a train and test set, with a test_size of 0.33 and respective random_state. Please use sklearn’s train_test_split for this: X_train, X_test, y_train, y_test = train_test_split(X_poly, y, test_size=0.33, random_state=seed). Please use X_poly from the last exercise.
- For each random seed train a ridge regressor with an alpha of your choice on the training set and calculate and print its MSE with respect to the test set
- Plot the predictions of the model for the whole range of x (not just X_train or X_test)

```
[ ]: from sklearn.model_selection import train_test_split
from sklearn.model_selection import cross_val_score
from sklearn.metrics import make_scorer
from sklearn.utils import shuffle
```

```

scorer = make_scorer(mean_squared_error)

plt.plot(x, y, "x")
plt.xlabel("x")
plt.ylabel("y")

# YOUR CODE HERE
raise NotImplementedError()

model = Ridge(alpha=10)
X_poly_, y_ = shuffle(X_poly, y)
scores = cross_val_score(model, X_poly_, y_, cv=10, scoring=scorer)
print("\n")
print(scores.mean())

```

0.4 Toy Problem

In this last part of today's assignment you can try out everything you've learned so far. In the next hidden cell a dataset is created that you should analyse and build a model for. Please use standard sklearn functions as the evaluation cell uses `model.predict(test_data)` to calculate the root mean square error.

Your goal is to build a model with a RMSE below 4, any model that can be evaluated with `model.predict` is fair game! Feel free to try different models, hyperparameter search, cross validation and other techniques to solve this exercise!

```

[ ]: # Hidden code generates X_train and y_train

from sklearn.datasets import load_boston
X, y = load_boston(return_X_y=True)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33,
    ↪random_state=711)
y_test = 5
X_test = 5

```

```

[ ]: # YOUR CODE HERE
raise NotImplementedError()

```

```

[ ]: # Hidden code uses model on test set and returns the RMSE.

print(rmse)
assert rmse < 4

```

```

[ ]:

```